

# Facial Recognition of Identical Twins using Convolutional Neural Networks

Springboard Capstone #2  
Anthony Johnson



# Outline

- Motivation
- Model
- Data
- Training
- Interpretation
- Conclusions

My brother



Me



# Motivation

- Facial recognition software has come a long way in the last few years, thanks to **convolutional neural networks (CNNs)**
  - iPhone X uses your face to unlock your phone
- How good exactly are CNNs?
- On a number of occasions I've been able to unlock my identical twin brother's iPhone with my face.
- Anything less than 100% accuracy is a major security concern

# Model

- **VGG-Face** model [1]
  - Open-source CNN
  - Pre-trained on hundreds of thousands of faces
- Use **transfer learning** to build on this

# Model Adaptation

- Define function

```
instantiate_model(num_dense, num_trainable)
```

- Creates same architecture as VGG-Face
- Load pre-trained weights, vgg\_face\_weights.h5
- Remove final layer
- Add a Dense layer with num\_dense = 100 nodes
- Add a Dense Activation('softmax') layer with 2 nodes
- Set all but the last num\_trainable = 5 layers as untrainable
- Input hyperparameters num\_dense and num\_trainable

# Model Architecture

- 145M parameters
  - 11M trainable
  - 134M untrainable

|     | Layer (type)             | Output Shape                     | Param #   |
|-----|--------------------------|----------------------------------|-----------|
| 1)  | zero_padding2d_131_input | (In (None, 224, 224, 3)          | 0         |
| 2)  | zero_padding2d_131       | (ZeroPadd (None, 226, 226, 3)    | 0         |
| 3)  | conv2d_161               | (Conv2D) (None, 224, 224, 64)    | 1792      |
| 4)  | zero_padding2d_132       | (ZeroPadd (None, 226, 226, 64)   | 0         |
| 5)  | conv2d_162               | (Conv2D) (None, 224, 224, 64)    | 36928     |
| 6)  | max_pooling2d_51         | (MaxPooling (None, 112, 112, 64) | 0         |
| 7)  | zero_padding2d_133       | (ZeroPadd (None, 114, 114, 64)   | 0         |
| 8)  | conv2d_163               | (Conv2D) (None, 112, 112, 128)   | 73856     |
| 9)  | zero_padding2d_134       | (ZeroPadd (None, 114, 114, 128)  | 0         |
| 10) | conv2d_164               | (Conv2D) (None, 112, 112, 128)   | 147584    |
| 11) | max_pooling2d_52         | (MaxPooling (None, 56, 56, 128)  | 0         |
| 12) | zero_padding2d_135       | (ZeroPadd (None, 58, 58, 128)    | 0         |
| 13) | conv2d_165               | (Conv2D) (None, 56, 56, 256)     | 295168    |
| 14) | zero_padding2d_136       | (ZeroPadd (None, 58, 58, 256)    | 0         |
| 15) | conv2d_166               | (Conv2D) (None, 56, 56, 256)     | 590080    |
| 16) | zero_padding2d_137       | (ZeroPadd (None, 58, 58, 256)    | 0         |
| 17) | conv2d_167               | (Conv2D) (None, 56, 56, 256)     | 590080    |
| 18) | max_pooling2d_53         | (MaxPooling (None, 28, 28, 256)  | 0         |
| 19) | zero_padding2d_138       | (ZeroPadd (None, 30, 30, 256)    | 0         |
| 20) | conv2d_168               | (Conv2D) (None, 28, 28, 512)     | 1180160   |
| 21) | zero_padding2d_139       | (ZeroPadd (None, 30, 30, 512)    | 0         |
| 22) | conv2d_169               | (Conv2D) (None, 28, 28, 512)     | 2359808   |
| 23) | zero_padding2d_140       | (ZeroPadd (None, 30, 30, 512)    | 0         |
| 24) | conv2d_170               | (Conv2D) (None, 28, 28, 512)     | 2359808   |
| 25) | max_pooling2d_54         | (MaxPooling (None, 14, 14, 512)  | 0         |
| 26) | zero_padding2d_141       | (ZeroPadd (None, 16, 16, 512)    | 0         |
| 27) | conv2d_171               | (Conv2D) (None, 14, 14, 512)     | 2359808   |
| 28) | zero_padding2d_142       | (ZeroPadd (None, 16, 16, 512)    | 0         |
| 29) | conv2d_172               | (Conv2D) (None, 14, 14, 512)     | 2359808   |
| 30) | zero_padding2d_143       | (ZeroPadd (None, 16, 16, 512)    | 0         |
| 31) | conv2d_173               | (Conv2D) (None, 14, 14, 512)     | 2359808   |
| 32) | max_pooling2d_55         | (MaxPooling (None, 7, 7, 512)    | 0         |
| 33) | conv2d_174               | (Conv2D) (None, 1, 1, 4096)      | 102764544 |
| 34) | dropout_21               | (Dropout) (None, 1, 1, 4096)     | 0         |
| 35) | conv2d_175               | (Conv2D) (None, 1, 1, 4096)      | 16781312  |
| 36) | dropout_22               | (Dropout) (None, 1, 1, 4096)     | 0         |
| 37) | conv2d_176               | (Conv2D) (None, 1, 1, 2622)      | 10742334  |
| 38) | flatten_11               | (Flatten) (None, 2622)           | 0         |
| 39) | dense_21                 | (Dense) (None, 100)              | 262300    |
| 40) | dense_22                 | (Dense) (None, 2)                | 202       |

Total params: 145,265,380

Trainable params: 11,004,836

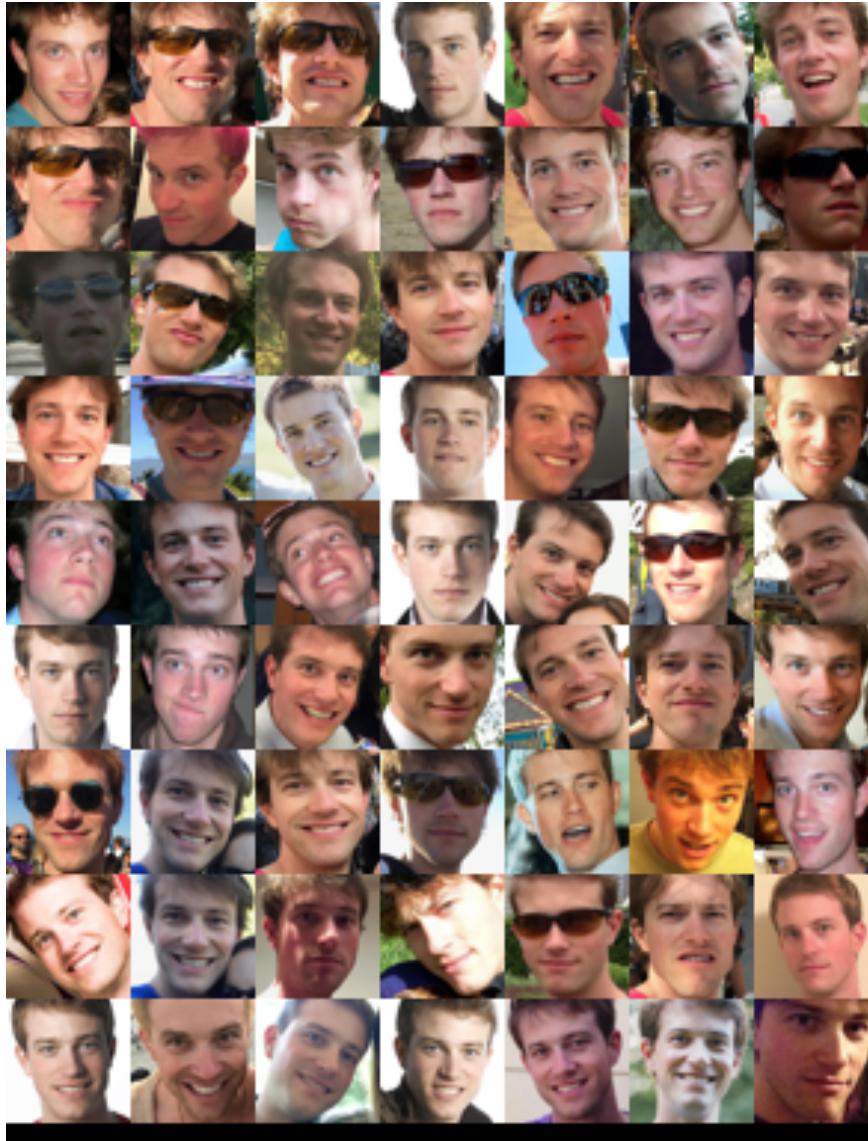
Non-trainable params: 134,260,544

# Data

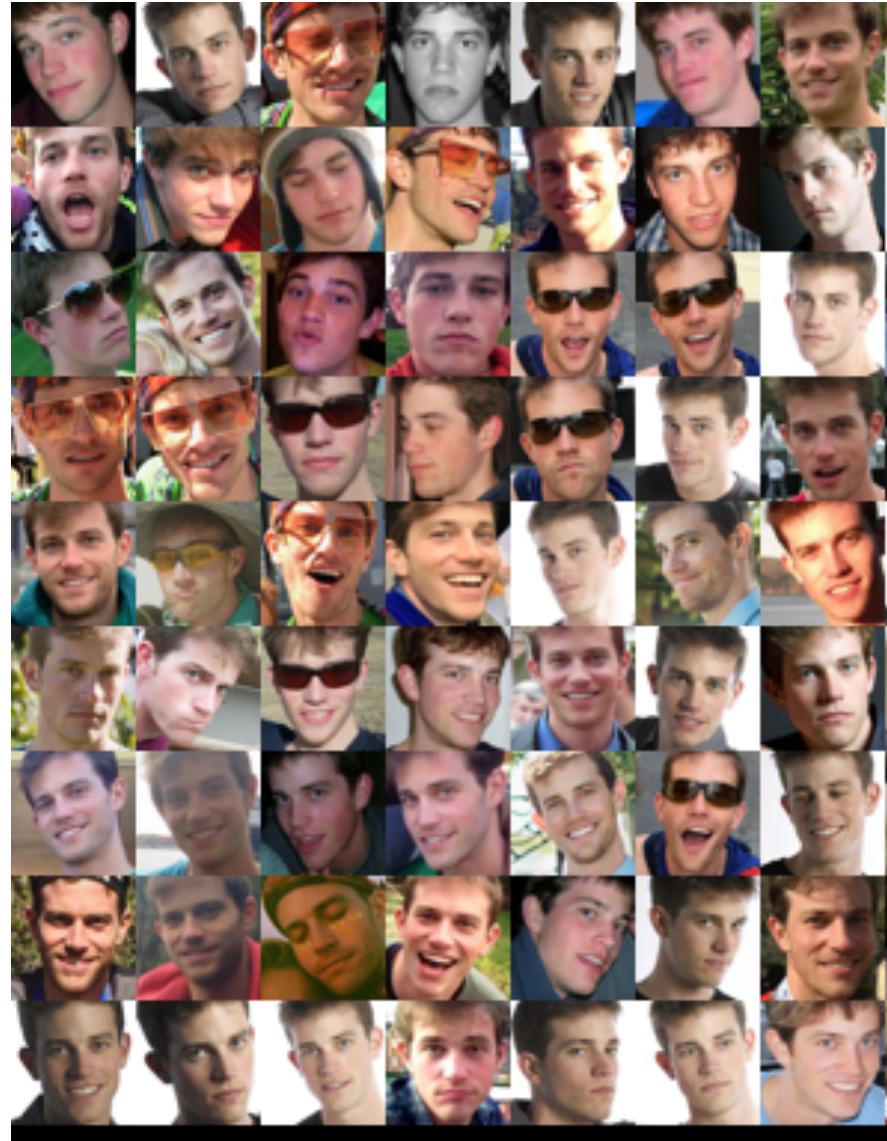
- I gathered digital photographs of me and my brother
- Square cropped each image
- Split into 3 datasets
  - Training = 118 images per twin
  - Validation = 27 images each per twin
  - Testing = 33 images each per twin
- Define **getXY()** function
  - Resizes images to 224x224
  - Stacks them together into trainX, valX, testX
  - Categorical dependent variable trainY, valY, testY

# Sample Training Images

Anthony



Paul

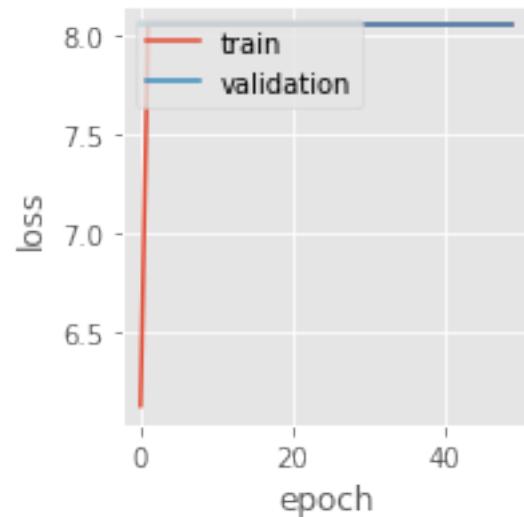
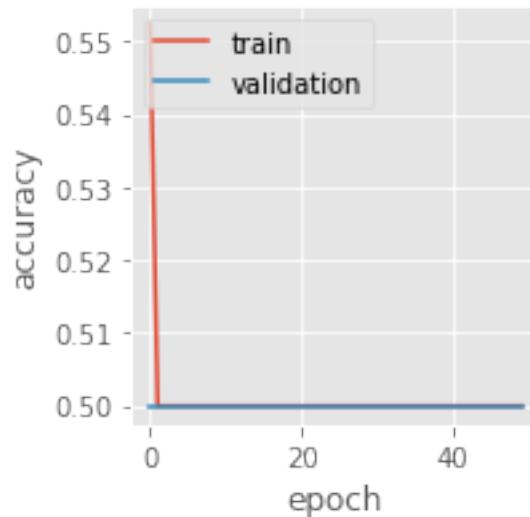


# Training

- Hyperparameters/settings to tune:
  - The number of layers that are trainable, **num\_trainable**
  - The learning rate, **lr**
  - The number of nodes, **num\_dense**, in the penultimate layer
  - The decay rate, **decay**, of the learning rate
  - The optimizer: SGD vs. Adam

# Training – Iteration #1

- Varied **num\_trainable** = [3, 4, 5, 6, 7, 8, 9]
- 120 nodes in the penultimate layer
- SGD optimizer with
  - Learning rate = 0.1 (Decay = 0)
  - Loss = ‘categorical\_crossentropy’
  - Metrics = ‘accuracy’
- These runs suffered from overfitting and thus did not converge.



# Training – Iteration #2

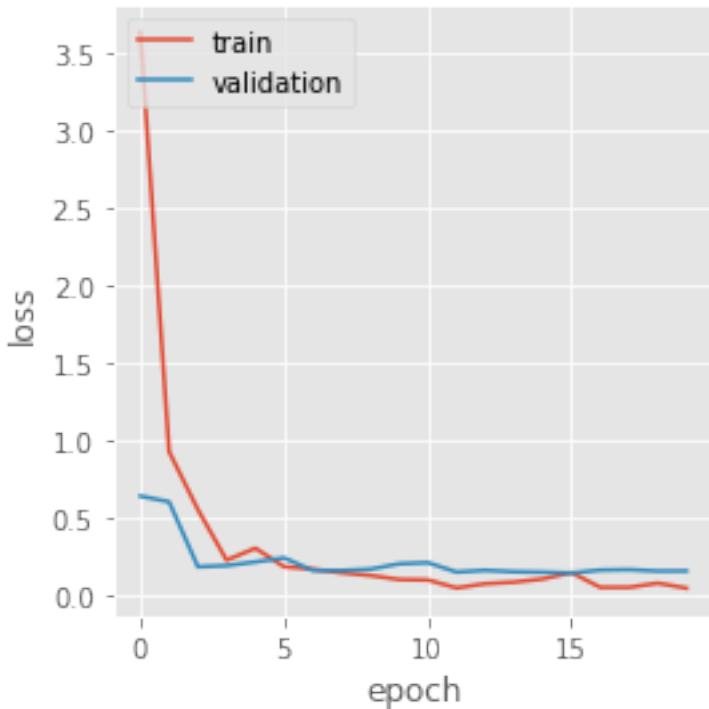
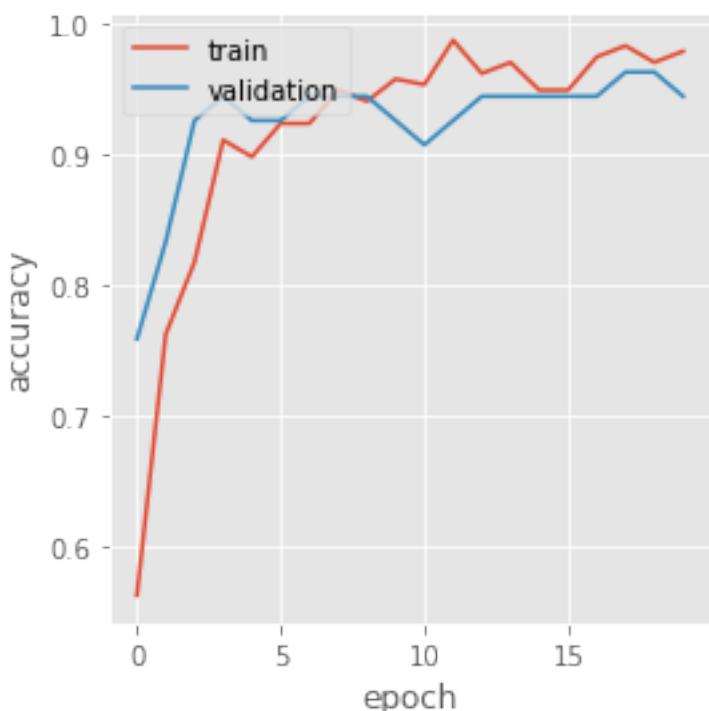
- Varied:
  - **num\_trainable** = [3, 4, 5, 6, 7, 8, 9]
  - Learning rate **lr** = [0.001, .005, .01, .05, .1, 1]
  - **num\_dense** = [10, 50, 100, 200]
- EarlyStopping condition
  - Patience of 4
  - Monitoring **val\_loss**
- These runs converged!
  - Training accuracies 90%-100%
  - Validation accuracies 85%-95%

# Training – Iteration #3

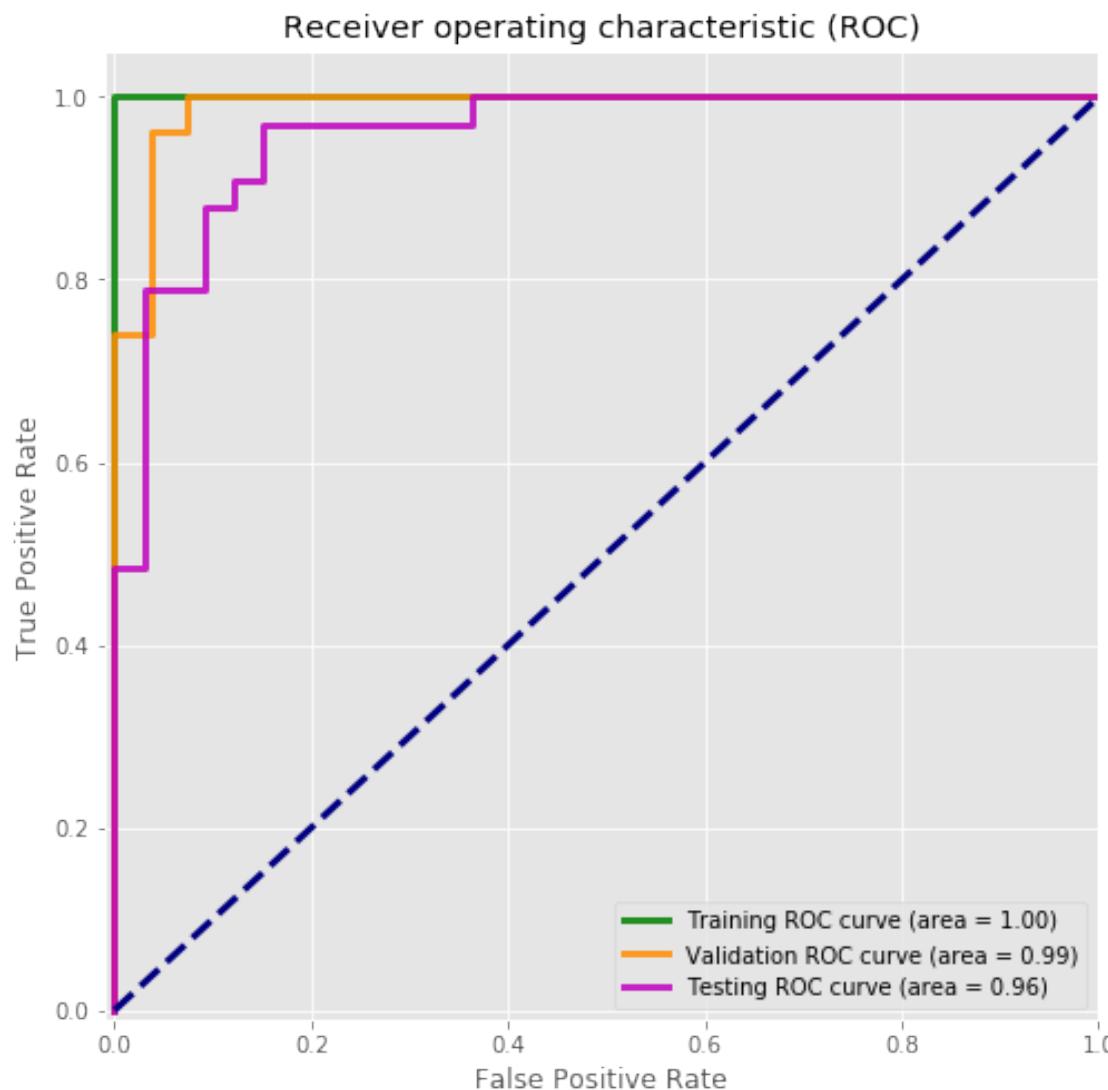
- Varied: Same as previous but varied decay
  - Decay = [.000005,.00001,.00002,.00004]
- EarlyStopping condition
  - Increased patience to 6
  - Monitoring **val\_loss**
- These runs also converged!
  - Training accuracies 90%-100%
  - Validation accuracies 85%-95%

# Best Model

- Parameters:
  - Num\_dense = 100
  - Num\_trainable = 5
  - Learning rate = 0.001
  - Decay = 0
  - Optimizer = SGD
  - EarlyStopping patience of 6 epochs
- Model accuracy:
  - Training: 100%
  - Validation: 94%
  - Testing: 88%



# Best Model – ROC Curve



# Best Model – Testing Results

| TESTING SET | PREDICTED - AJ | PREDICTED - PJ |
|-------------|----------------|----------------|
| ACTUAL - AJ | 29             | 4              |
| ACTUAL - PJ | 4              | 29             |

Actual = AJ, Predicted = PJ



Potential causes of these misclassifications:

- Bad Lighting
- Sunglasses
- Unusual facial tilt or expressions

Actual = PJ, Predicted = AJ



# Interpretation

- **LIME** = Local Interpretable Model-Agnostic Explanations [2]
  - *LimeImageExplainer()* identifies features in an image that increase the probability (green) or decrease the probability (red) of a certain class
  - Red patches indicate a higher probability of Paul and green patches indicate a higher probability of Anthony

Paul

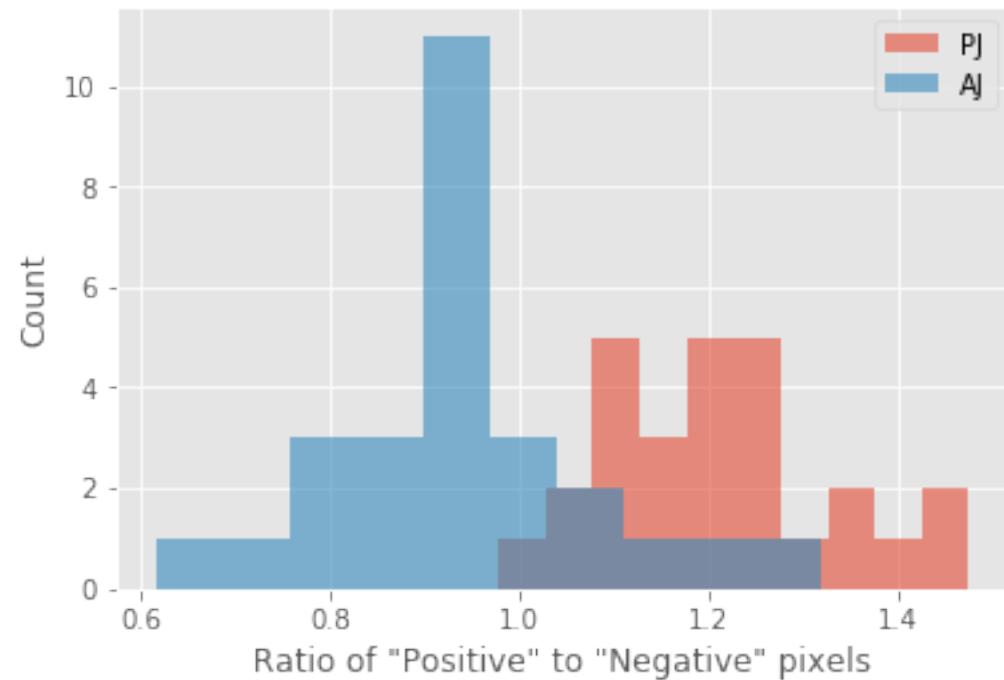


Anthony



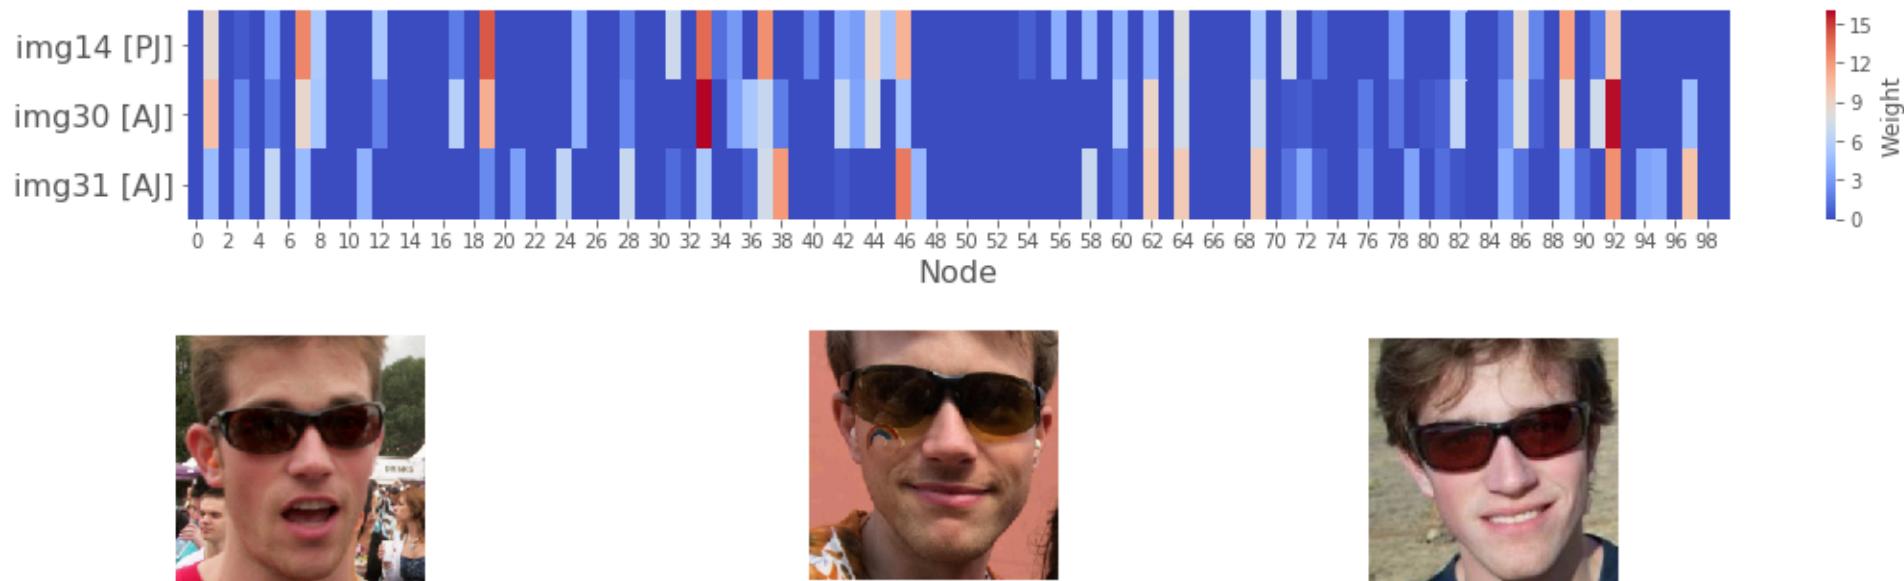
# Interpretation

- Analysis:
  - Sum the total number of green and red pixels for both AJ and PJ
  - Divide total “positive” pixels by total “negative” pixels
  - Create histogram (right)
- In PJ images, the ratio of red to green pixels was between 1.0 and 1.5, averaged to 1.21.
- In AJ images, the equivalent ratio varied much more widely, between 0.6 and 1.3.



# Direct Face Comparison

- Penultimate layer node weights
  - Middle row is misclassified



Img14 [PJ]



Img30 [AJ]  
[thinks it is PJ]



Img31 [AJ]

# Conclusions

- The final convolutional neural network:
  - Optimized the learning rate (.001), number of trainable layers (5), and number of nodes in the final layers (100)
  - Converged in <5 epochs
  - Final testing accuracy = 88%
  - Final testing ROC AUC = 0.96
- Overall, the model failed in cases with bad or inconsistent lighting or sometimes when we were wearing sunglasses
- Model seems to have slightly easier time recognizing PJ than AJ
  - Maybe PJ images are "better" or more unique by chance

# Future Work

- While decent for identical twins, this model would be unacceptable in a security application
- To improve the model:
  - Train more images (the more the better)
  - Further optimize the hyperparameters
  - Optimize the optimizer
- I would like to dig deeper into interpreting and visualizing CNNs
  - With millions of parameters, CNNs are difficult to interpret
  - But this is still a new field, and anything is possible